

The Tell: Which Pitch-Sequencing Giveaways Survive New Games, and Why Hitters Leave Them on the Table

Introduction. Advance scouts sell sequencing tells (“after a slider he doubles up”), and clubs act on them in hitter meetings and pitcher development. No public method separates a real tell from a pattern found by searching hundreds of pitchers, counts and previous pitches. We ask which pitchers give away the next pitch beyond count and batter hand, how much of a scouted tell survives into new games, whether hitters punish it, and whether the 2026 automated ball-strike (ABS) system changed it.

Methods. We use all 2,049,352 regular-season Statcast pitches from 2024–2026 (1,793 pitcher-seasons). For each pitcher-season we fit nested predictive models of pitch type (batter hand, then count, then the previous pitch) with hierarchical Dirichlet smoothing, scored leave-one-game-out. The out-of-sample gain from the previous pitch is TELL, in bits per pitch. A concrete tell is the count, hand and previous-pitch situation that shifts one pitch type most. We select it on April–June or on 2025 and score it only on later games. Permutation e-values, multiplied game by game, form an in-season alarm whose chance of ever falsely firing is at most 5%. Cost regressions use pitcher-season and pitch-type-by-count fixed effects. ABS is evaluated on paired pitchers against a 2024–25 placebo.

Results.


- Batter hand and count carry 102 and 58 millibits about the next pitch; the previous pitch adds 8.1. A hitter guessing the likeliest pitch gains 0.6 points of accuracy on average, and 5 against Dylan Cease.
- Selection inflates tells. Each pitcher’s largest April–June tell (26 points) keeps 10 points in July–September, in the same direction 78% of the time (Figure 1B). From 2025 to 2026: 28 to 13 points, 79%. Doubling-up tells hold best (14 versus 8 points within 2025).
- The ten largest 2025 tells with enough 2026 data all kept their direction (Table 1).
- TELL is reliable within a season (split-half 0.73) but not across seasons ($r = 0.19$, versus 0.62 for count usage).
- Hitters do not cash in. Per bit of predictability, pitches drew 1.5 more points of whiffs ($t = 4.7$) and 0.59 more runs saved per 100 pitches ($t = 5.3$) (Figure 1C).
- The alarm fired for 24% of 2025 pitchers, at a median of 290 pitches (Figure 1D).
- ABS lowered the top of the called zone by 1.2 inches (placebo 0.2) and moved TELL by -0.04 millibits (95% CI -1.47 to $+1.39$).

Conclusion. Sequencing tells are real, season-specific, and less than half as large in new games as when found. Scouts should discount them accordingly. Breaking them is not worth development time at observed levels. Both can use an alarm that flags a detectable pattern with a guaranteed false-alarm rate. Code: github.com/aidaneichman/the-tell.


Table 1. Largest 2025 tells (50+ cue pitches; 25+ in 2026), re-scored on 2026. Next-pitch share after the cue versus otherwise (%).

Pitcher	Situation	Next pitch	2025	2026
Dylan Lee	0-1 vs RHB, after slider	slider	75 vs 9	69 vs 25
Jack Leiter	1-2 vs LHB, after 4-seam	4-seam	22 vs 68	33 vs 50
Kevin Gausman	0-2 vs LHB, after 4-seam	4-seam	60 vs 15	51 vs 26
Chris Sale	0-1 vs RHB, after slider	4-seam	22 vs 66	30 vs 55
Gregory Soto	0-1 vs RHB, after sinker	sinker	62 vs 19	46 vs 24
Matthew Boyd	2-2 vs RHB, after change	4-seam	71 vs 30	56 vs 28
Hunter Brown	0-1 vs RHB, after sinker	sinker	62 vs 22	64 vs 28
Kris Bubic	1-0 vs RHB, after 4-seam	4-seam	54 vs 16	67 vs 24
Drew Rasmussen	1-2 vs LHB, after cutter	cutter	7 vs 45	4 vs 30
Dustin May	1-2 vs LHB, after sweeper	sweeper	31 vs 68	19 vs 42


A One tell: Yusei Kikuchi, 1–0 count, right-handed batter



B Each pitcher's biggest tell, re-scored on new games (n = 110)



C Hitters do not cash in: predictable pitches miss more bats



D An in-season alarm with a false-alarm guarantee




Figure 1. (C) Decile bins, 95% intervals. (D) Dots mark when each alarm fired.