

The Tell: Which Pitch-Sequencing Giveaways Survive New Games

Introduction. Advance reports flag sequencing tells (“after a slider he doubles up”), and clubs act on them. Public sequencing work has not tested whether pitcher-specific cues survive new games. We ask how much the previous pitch gives away beyond count and batter hand, how much cues survive, whether predictability costs pitchers, and whether it changed in 2026, the first automated ball-strike (ABS) challenge season.

Methods. 2,049,352 Statcast pitches, 2024–2026 (1,793 pitcher-seasons with 250+ pitches). Nested models predict pitch type from batter hand, then count, then previous pitch, on unseen games. TELL is the previous-pitch gain per pitcher-season, over pitches after a plate appearance’s first. A cue is the count, hand and previous pitch that most shifts one pitch type’s share. Each pitcher’s top cue is found in April–June (or 2025), scored only on later games, and also net of the league. Outcome regressions fix pitcher-season and pitch type $\times$ count $\times$ hand, controlling velocity; a second adds the previous pitch’s result. An in-season alarm multiplies game-by-game permutation e-values within count and batter hand. 2025–26 changes are compared with 2024–25, with TELL paired by pitcher.

Results.

- In 2025, hand and count carry 102 and 58 millibits; the previous pitch adds 12.8. Using the previous pitch, guessing the likeliest pitch is right 0.8 points more often, 7 against Dylan Cease.
- The top April–June cue shifts the next-pitch share by 26 points when found and 10 on July–September games (78% same direction; no tell: 0, 50%; Figure 1B). Net of the league, 7 points (70%).
- The 10 largest 2025 cues all kept direction in 2026, shrinking 44 to 29 points (Table 1).
- TELL’s within-season split-half reliability is 0.75; across seasons, $r = 0.22$ (count gain: 0.62).
- Pitch-level predictability is not associated with worse outcomes. Within pitcher-season, from its 10th to 90th percentile, whiffs per swing rise 1.1 points ($t = 6.8$) and pitchers save 0.26 more runs per 100 pitches ($t = 4.7$); controlling the previous pitch’s result, 0.9 and 0.24.
- The alarm fired for 25% of 2025 pitchers (median 436 pitches), 1% of cue-free controls, and 1.6% of simulations with the count responding to pitches (Figure 1C).
- In 2026 the called zone’s top was 1.2 inches lower (2024→25: 0.2); TELL changed +0.01 millibits (95% CI -1.86 to $+1.88$; 2024→25: -0.49). The zone moved; sequencing did not detectably.


Conclusion. Sequencing tells are real but shrink by a third to three-quarters when held out, and overall predictability is far less stable across seasons than within one. Advance scouts should validate cues on held-out games, discount their size, and use the alarm to trigger review.

Table 1. Top 2025 cues (50+ pitches; 25+ in 2026). Shift: share of the pitch after the cue minus after other previous pitches, same count and hand.

Pitcher	Situation	Next pitch	Found 2025	Held out 2026	<i>n</i> 2026
Dylan Lee	0-1 vs RHB, after slider	slider	+66	+44	62
Jack Leiter	1-2 vs LHB, after 4-seam	4-seam	-46	-17	39
Kevin Gausman	0-2 vs LHB, after 4-seam	4-seam	+45	+25	74
Chris Sale	0-1 vs RHB, after slider	4-seam	-44	-25	111
Gregory Soto	0-1 vs RHB, after sinker	sinker	+43	+22	48
Matthew Boyd	2-2 vs RHB, after change	4-seam	+41	+28	54
Hunter Brown	0-1 vs RHB, after sinker	sinker	+40	+36	50
Kris Bubic	1-0 vs RHB, after 4-seam	4-seam	+38	+43	36
Drew Rasmussen	1-2 vs LHB, after cutter	cutter	-38	-26	51
Dustin May	1-2 vs LHB, after sweeper	sweeper	-37	-23	31


A A cue that held up: Kikuchi, 2025, 1-0 vs RHB

next-pitch mix after a slider vs after anything else



B Every pitcher's top cue, held out (*n* = 110)

found 26 pts, held out 10; same direction 78% (no tell: 0, 50%)



C An in-season alarm

permutes pitch order within count and batter hand, game by game




Figure 1. (A) Apr-Jun 2025 vs Jul-Sep. (B) Top Apr-Jun cues (30+ pitches each way), re-scored Jul-Sep (10+). (C) Alarm paths; faded after firing.